

# Top 30

## Kubernetes Interview Questions & Answers

Asked at Top MNCs (Google, Amazon, Microsoft, Meta, Apple, Netflix)

### SECTION 1: FUNDAMENTALS & CORE CONCEPTS

#### Q1: What is Kubernetes and why is it so popular?

##### Answer:

Kubernetes (K8s) is an open-source container orchestration platform originally developed by Google and now maintained by the Cloud Native Computing Foundation (CNCF). It automates the deployment, scaling, and management of containerized applications.

##### Key Features:

- **Container Orchestration:** Manages containerized applications across clusters of machines
- **Self-Healing:** Automatically restarts failed containers, replaces unresponsive nodes, and reschedules workloads
- **Auto-Scaling:** Scales applications horizontally (more pods) or vertically (more resources) based on demand
- **Load Balancing:** Distributes traffic efficiently between pods
- **Rolling Updates & Rollbacks:** Allows zero-downtime deployments with automatic rollback capability
- **Declarative Configuration:** Infrastructure as Code using YAML manifests

##### Popularity Reasons:

- Over 70% of Fortune 500 companies use Kubernetes in production
- Cloud-agnostic (works with AWS, Azure, GCP, on-premises)

- Simplifies microservices architecture management
- Strong ecosystem and community support
- Industry standard for container orchestration (replaced Docker Swarm)

## **Q2: Explain the Kubernetes Architecture**

### **Answer:**

Kubernetes follows a master-worker (control plane-data plane) architecture:

#### **Control Plane Components (Master Node):**

- 1. API Server (`kube-apiserver`)**
  - Front-end of Kubernetes cluster
  - Handles all REST API requests
  - Validates and persists data to etcd
  - Acts as the gateway for cluster communication
- 2. etcd**
  - Distributed key-value store
  - Stores entire cluster state and configuration
  - Single source of truth for cluster data
  - Requires backup strategy for disaster recovery
- 3. Scheduler (`kube-scheduler`)**
  - Assigns newly created pods to nodes
  - Considers resource availability, constraints, and policies
  - Makes scheduling decisions based on pod requirements
- 4. Controller Manager (`kube-controller-manager`)**
  - Runs controller processes
  - Manages ReplicaSets, Deployments, StatefulSets, DaemonSets
  - Handles node failures and maintains desired state

#### **Worker Node Components:**

- 1. Kubelet**
  - Agent running on each worker node
  - Ensures containers run as specified
  - Reports node and pod status to API server
  - Executes pod lifecycle management
- 2. Kube-proxy**
  - Network proxy on each node
  - Manages network rules for pod communication
  - Implements service abstraction and load balancing
  - Maintains network connectivity
- 3. Container Runtime**

- Executes containers (Docker, containerd, CRI-O)
- Pulls images and runs containers
- Manages container lifecycle

### **How They Work Together:**

User submits deployment → API Server validates → Stored in etcd → Scheduler assigns pods to nodes → Kubelet creates containers → Kube-proxy manages networking

## **Q3: What is a Pod and how does it differ from Docker containers?**

### **Answer:**

A **Pod** is the smallest deployable unit in Kubernetes, not containers themselves.

### **Key Pod Characteristics:**

- Can contain one or more containers (usually one)
- Containers in a pod share networking namespace (same IP address)
- Containers share storage volumes
- Containers can communicate via localhost
- All containers in a pod are created/destroyed together
- Ephemeral by nature (can be created/destroyed anytime)

### **Pod vs Container:**

<b>Aspect</b>	<b>Pod</b>	<b>Container</b>
<b>Creation</b>	Kubernetes manages	Docker manages
<b>Networking</b>	Shared IP	Individual IP per container
<b>Lifecycle</b>	Managed by controllers	Managed by runtime
<b>Storage</b>	Can share volumes	Isolated storage
<b>Scaling</b>	Via ReplicaSets/Deployments	Manual scaling

### **Example Pod YAML:**

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - name: nginx
    image: nginx:1.21
    ports:
    - containerPort: 80
  resources:
    requests:
      cpu: 100m
      memory: 128Mi
    limits:
      cpu: 500m
      memory: 512Mi
```

### **When to use Multi-Container Pods:**

- Sidecar containers (logging, monitoring)
- Ambassador containers (proxy, configuration management)
- Init containers (setup tasks)
- Tightly coupled components that must share resources

## **Q4: What is a Node in Kubernetes?**

### **Answer:**

A **Node** is a worker machine in the Kubernetes cluster that runs containerized applications.

### **Node Types:**

1. **Worker Node:** Runs application pods
2. **Master/Control Plane Node:** Manages cluster state and scheduling

### **Node Characteristics:**

- Can be physical machine or virtual machine
- Registers itself with API server
- Has unique name and IP address
- Runs kubelet, kube-proxy, and container runtime
- Has labels for scheduling decisions

### **Node Status Contains:**

- **Address:** Hostname, internal IP, external IP
- **Condition:** Ready, MemoryPressure, DiskPressure, PIDPressure, NetworkUnavailable
- **Capacity:** Total resources available (CPU, memory, pods)
- **Allocatable:** Resources available for pod scheduling

#### **Check Node Status:**

```
kubectl get nodes  
kubectl describe node <node-name>  
kubectl top nodes # Resource usage
```

## **Q5: What is a Namespace and why use them?**

#### **Answer:**

A **Namespace** is a virtual cluster within a physical Kubernetes cluster, providing logical isolation.

#### **Key Benefits:**

- **Resource Isolation:** Separate resources for different teams/projects
- **Multi-Tenancy:** Multiple teams can share same cluster safely
- **Resource Quotas:** Limit resources per namespace
- **RBAC:** Different access controls per namespace
- **Network Policies:** Isolate network traffic between namespaces
- **Naming Scope:** Same names can exist in different namespaces

#### **Default Namespaces:**

- default: Default namespace for resources
- kube-system: Kubernetes system components
- kube-public: Public resources
- kube-node-lease: Node heartbeat data

#### **Create and Use Namespace:**

```
kubectl create namespace dev  
kubectl run nginx --image=nginx --namespace=dev  
kubectl get pods --namespace=dev
```

#### **Use Cases:**

- Separate development, staging, production environments
- Multi-team environments with resource isolation

- Compliance requirements for data separation

## **Q6: What are Labels and Selectors?**

### **Answer:**

**Labels** are key-value pairs attached to Kubernetes objects for organization and identification.

**Selectors** filter resources based on labels.

### **Example:**

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
  labels:
    app: web-server
    environment: production
    tier: frontend
spec:
  containers:
    - name: nginx
      image: nginx
```

### **Selector Types:**

#### **1. Equality-based:**

```
kubectl get pods -l environment=production
kubectl get pods -l app!=web-server
```

#### **2. Set-based:**

```
kubectl get pods -l "environment in (prod, staging)"
kubectl get pods -l "tier notin (backend, database)"
```

### **Use Cases:**

- Service discovery and load balancing
- Pod scheduling (node selectors)
- RBAC policies
- Network policies
- Resource quotas
- Monitoring and logging

# SECTION 2: DEPLOYMENTS & WORKLOAD MANAGEMENT

## Q7: Explain Deployments and their lifecycle

**Answer:**

A **Deployment** is a higher-level abstraction that manages Pods and ReplicaSets, providing declarative updates.

### Deployment Structure:

- Deployment → ReplicaSet (v1) → Pod Replicas (old version)
- Deployment (after update) → ReplicaSet (v2) → Pod Replicas (new version)

### Key Features:

- **Declarative Updates:** Define desired state, Kubernetes handles updates
- **Rollback Support:** Easy rollback to previous versions
- **Rolling Updates:** Gradually replace old pods with new ones
- **Scaling:** Horizontal pod scaling
- **Self-Healing:** Replaces failed pods

### Deployment YAML Example:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 1
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
  spec:
    containers:
```

```

- name: nginx
  image: nginx:1.21
  ports:
  - containerPort: 80
  resources:
    requests:
      cpu: 100m
      memory: 128Mi
    limits:
      cpu: 500m
      memory: 512Mi

```

### **Deployment Lifecycle:**

1. **Creation:** Deployment creates ReplicaSet, which creates Pods
2. **Scaling:** Update replicas field to scale up/down
3. **Update:** Change image/config triggers rolling update
4. **Rollback:** Revert to previous ReplicaSet if issues occur
5. **Deletion:** Cascading deletion of ReplicaSets and Pods

### **Common Commands:**

```

kubectl create deployment nginx --image=nginx
kubectl scale deployment nginx --replicas=5
kubectl set image deployment/nginx nginx=nginx:1.22
kubectl rollout status deployment/nginx
kubectl rollout history deployment/nginx
kubectl rollout undo deployment/nginx # Rollback to previous

```

## **Q8: What is the difference between Deployment, StatefulSet, and DaemonSet?**

### **Answer:**

<b>Feature</b>	<b>Deployment</b>	<b>StatefulSet</b>	<b>DaemonSet</b>
<b>Use Case</b>	Stateless apps	Stateful apps (databases)	Node-level services
<b>Pod Identity</b>	Interchangeable	Unique, stable (pod-0, pod-1)	One per node

<b>Storage</b>	Shared/No persistence	Persistent per pod	Usually none
<b>Ordering</b>	Parallel	Ordered creation/deletion	Parallel
<b>Scaling</b>	Easy, any replica	Ordered, carefully	One per node
<b>Examples</b>	Web servers, APIs	MySQL, PostgreSQL, Kafka	Logging agents, monitoring
<b>Network Identity</b>	ClusterIP	Stable, predictable	Variable
<b>Update Strategy</b>	Rolling Update, Recreate	Rolling with ordering	Rolling on each node

### **StatefulSet Example:**

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: mysql
spec:
  serviceName: mysql
  replicas: 3
  selector:
    matchLabels:
      app: mysql
  template:
    metadata:
      labels:
        app: mysql
  spec:
    containers:
      - name: mysql
        image: mysql:8.0
        ports:
          - containerPort: 3306
    volumeMounts:
      - name: mysql-storage

```

```
        mountPath: /var/lib/mysql
volumeClaimTemplates:
- metadata:
  name: mysql-storage
spec:
  accessModes: ["ReadWriteOnce"]
  resources:
    requests:
      storage: 10Gi
```

### **DaemonSet Example:**

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: fluentd-logging
spec:
  selector:
    matchLabels:
      name: fluentd-logging
  template:
    metadata:
      labels:
        name: fluentd-logging
    spec:
      tolerations:
        - key: node-role.kubernetes.io/master
          effect: NoSchedule
      containers:
        - name: fluentd
          image: fluentd:v1.14
          volumeMounts:
            - name: varlog
              mountPath: /var/log
            - name: varlibdockercontainers
              mountPath: /var/lib/docker/containers
              readOnly: true
      volumes:
        - name: varlog
          hostPath:
            path: /var/log
        - name: varlibdockercontainers
```

```
hostPath:  
  path: /var/lib/docker/containers
```

### When to Use:

- **Deployment:** Web apps, microservices, REST APIs, stateless workloads
- **StatefulSet:** Databases, message queues, distributed systems (Elasticsearch, Kafka)
- **DaemonSet:** Logging agents (Fluentd), monitoring (Prometheus Node Exporter), network plugins

## Q9: What is a ReplicaSet?

### Answer:

A **ReplicaSet** ensures a specified number of pod replicas are running at all times.

### Key Characteristics:

- Created automatically by Deployments (not used directly)
- Continuously monitors pods and creates/deletes as needed
- Uses selectors to identify pods
- Supports both equality and set-based selectors

### ReplicaSet vs Replication Controller:

- **ReplicaSet:** Newer, supports set-based selectors
- **Replication Controller:** Older, only equality-based selectors

### ReplicaSet YAML:

```
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:  
  name: nginx-rs  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: nginx  
  template:  
    metadata:  
      labels:  
        app: nginx  
    spec:  
      containers:
```

```
- name: nginx
  image: nginx:1.21
```

## SECTION 3: SERVICES & NETWORKING

### Q10: What is a Kubernetes Service and why is it needed?

#### Answer:

A **Service** abstracts a set of pods and provides a stable network endpoint for accessing them.

#### Why Services are Needed:

- Pods are ephemeral (IPs change when pods restart)
- Services provide stable IP and DNS name
- Enable service discovery within cluster
- Load balance traffic across pods
- Expose services internally or externally

#### Service YAML Example:

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: nginx
  ports:
    - protocol: TCP
      port: 80      # Service port
      targetPort: 8080 # Pod port
  type: ClusterIP
```

### Q11: Explain the four types of Kubernetes Services

#### Answer:

#### 1. ClusterIP (Default)

- Internal service accessible only within cluster
- Assigns virtual IP from cluster IP range
- Used for internal pod-to-pod communication
- No external access

```
apiVersion: v1
kind: Service
metadata:
  name: internal-service
spec:
  type: ClusterIP
  selector:
    app: api
  ports:
    - port: 80
      targetPort: 8080
```

## 2. **NodePort**

- Exposes service on each node's IP at a static port (30000-32767)
- Accessible from outside cluster via <NodeIP>:<NodePort>
- Creates ClusterIP service internally
- Less secure, suitable for development

```
apiVersion: v1
kind: Service
metadata:
  name: web-service
spec:
  type: NodePort
  selector:
    app: web
  ports:
    - port: 80
      targetPort: 8080
      nodePort: 30080 # Optional, auto-assigned if not specified
```

## 3. **LoadBalancer**

- Exposes service via cloud provider's load balancer
- Gets external IP from cloud provider
- Each service creates separate load balancer (expensive)
- Recommended for production external services

```
apiVersion: v1
kind: Service
metadata:
  name: api-lb
spec:
  type: LoadBalancer
  selector:
    app: api
  ports:
    - port: 80
      targetPort: 8080
```

#### 4. **ExternalName**

- Maps service to external hostname/domain
- Returns CNAME record pointing to external service
- No proxy, direct DNS redirect

```
apiVersion: v1
kind: Service
metadata:
  name: external-db
spec:
  type: ExternalName
  externalName: mysql.example.com
```

#### **Service Selection Summary:**

- Internal communication → **ClusterIP**
- Development/Testing → **NodePort**
- Production external → **LoadBalancer**
- External resource mapping → **ExternalName**

## **Q12: What is an Ingress and how does it differ from Service?**

#### **Answer:**

An **Ingress** manages external HTTP/HTTPS access to services within a cluster.

#### **Ingress vs Service:**

<b>Aspect</b>	<b>Service</b>	<b>Ingress</b>
<b>Layer</b>	Layer 4 (Transport)	Layer 7 (Application)
<b>Protocol</b>	TCP/UDP	HTTP/HTTPS
<b>Routing</b>	Port-based	Path/hostname-based
<b>Use Case</b>	Direct pod access	Web application routing
<b>Cost</b>	Per service load balancer	Single load balancer for multiple apps

#### **Ingress YAML Example:**

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: multi-app-ingress
  annotations:
    cert-manager.io/cluster-issuer: letsencrypt-prod
spec:
  ingressClassName: nginx
  tls:
    - hosts:
        - app1.example.com
        - app2.example.com
      secretName: tls-secret
  rules:
    - host: app1.example.com
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: app1-service
                port:
                  number: 80
    - host: app2.example.com
      http:
        paths:
          - path: /api
            pathType: Prefix
            backend:
              service:
                name: app2-service
                port:
                  number: 8080
          - path: /static
            pathType: Prefix
            backend:
              service:
                name: static-service
                port:
                  number: 3000
```

### **Ingress Requirements:**

- Ingress Controller (nginx, Traefik, HAProxy)
- Service for each backend application
- DNS pointing to ingress IP

## **Q13: What is a Network Policy?**

### **Answer:**

A **Network Policy** defines rules for pod-to-pod and pod-to-external communication.

### **Default Behavior:**

- By default, all pods can communicate with all pods
- Network policies restrict this communication

### **Network Policy YAML:**

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: app-network-policy
spec:
  podSelector:
    matchLabels:
      role: database
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              role: app
      ports:
        - protocol: TCP
          port: 3306
  egress:
    - to:
        - podSelector:
            matchLabels:
              role: app
      ports:
        - protocol: TCP
```

port: 8080

### **Key Concepts:**

- podSelector: Which pods the policy applies to
- policyTypes: Ingress (incoming) and/or Egress (outgoing) rules
- from/to: Source/destination for traffic
- ports: Protocol and port restrictions

### **Best Practices:**

- Implement least privilege access
- Deny all by default, allow specific traffic
- Use pod labels for fine-grained control
- Monitor and audit network policies

## **SECTION 4: STORAGE & PERSISTENCE**

### **Q14: Explain Persistent Volumes (PV) and Persistent Volume Claims (PVC)**

#### **Answer:**

Storage in Kubernetes is managed via Persistent Volumes and claims.

#### **Persistent Volume (PV):**

- Storage resource provisioned by administrator
- Cluster-level resource (not namespaced)
- Has lifecycle independent of pods
- Supports various storage backends (NFS, iSCSI, cloud storage)

#### **Persistent Volume Claim (PVC):**

- User's request for storage
- Binds to a PV
- Can specify storage size and access modes
- Namespaced resource

#### **PV and PVC YAML:**

```
# Persistent Volume
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mysql-pv
spec:
```

```
capacity:
  storage: 50Gi
accessModes:
- ReadWriteOnce
persistentVolumeReclaimPolicy: Retain
storageClassName: fast-ssd
nfs:
  server: nfs.example.com
  path: "/data/mysql"
```

```
# Persistent Volume Claim
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pvc
  namespace: production
spec:
  accessModes:
- ReadWriteOnce
  storageClassName: fast-ssd
  resources:
    requests:
      storage: 50Gi
```

```
# Pod using PVC
apiVersion: v1
kind: Pod
metadata:
  name: mysql
spec:
  containers:
- name: mysql
  image: mysql:8.0
  volumeMounts:
- name: mysql-storage
  mountPath: /var/lib/mysql
  volumes:
- name: mysql-storage
  persistentVolumeClaim:
    claimName: mysql-pvc
```

### **Access Modes:**

- ReadWriteOnce (RWO): Single node read-write
- ReadOnlyMany (ROX): Multiple nodes read-only
- ReadWriteMany (RWX): Multiple nodes read-write

### **Reclaim Policies:**

- Retain: Manual cleanup required
- Delete: Automatically deleted with PVC
- Recycle: Basic scrub (deprecated)

## **Q15: What is a StorageClass?**

### **Answer:**

A **StorageClass** enables dynamic provisioning of Persistent Volumes.

### **Benefits:**

- Automatically create PVs based on PVC requests
- No manual PV creation needed
- Different storage tiers/types
- Customizable provisioning parameters

### **StorageClass Example:**

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: fast-ssd
provisioner: kubernetes.io/aws-ebs
allowVolumeExpansion: true
parameters:
  type: gp3
  iops: "3000"
  throughput: "125"
  encrypted: "true"
volumeBindingMode: WaitForFirstConsumer
```

### **Using StorageClass:**

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
```

```
name: app-storage
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: fast-ssd
  resources:
    requests:
      storage: 100Gi
```

### **Dynamic Provisioning Flow:**

1. User creates PVC with storageClassName
2. Controller detects PVC creation
3. Provisioner creates PV based on StorageClass parameters
4. PVC binds to newly created PV
5. Pod can use the storage

## **SECTION 5: CONFIGURATION & SECRETS**

### **Q16: What are ConfigMaps and Secrets?**

#### **Answer:**

#### **ConfigMap:**

- Stores non-sensitive configuration data
- Key-value pairs or files
- Mounted as files or environment variables
- Data visible in plain text

#### **Secret:**

- Stores sensitive information (passwords, API keys)
- Base64 encoded (not encrypted by default)
- Automatically injected via volumes or environment variables
- Encryption at rest can be enabled

#### **ConfigMap Example:**

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  DATABASE_HOST: postgres.prod.svc.cluster.local
  DATABASE_PORT: "5432"
```

```
LOG_LEVEL: info
app.conf: |
  server {
    listen 8080;
    location / {
      proxy_pass http://backend;
    }
  }
```

### **Using ConfigMap:**

```
apiVersion: v1
kind: Pod
metadata:
  name: app-pod
spec:
  containers:
    - name: app
      image: myapp:1.0
      envFrom:
        - configMapRef:
            name: app-config
  volumeMounts:
    - name: config-volume
      mountPath: /etc/config
  volumes:
    - name: config-volume
      configMap:
        name: app-config
```

### **Secret Example:**

```
apiVersion: v1
kind: Secret
metadata:
  name: db-credentials
type: Opaque
data:
  username: dXNlcmt5hbWU= # base64 encoded "username"
  password: cGFzc3dvcmQ= # base64 encoded "password"
```

### **Using Secret:**

```
apiVersion: v1
kind: Pod
metadata:
  name: db-pod
spec:
  containers:
    - name: mysql
      image: mysql:8.0
      env:
        - name: MYSQL_ROOT_PASSWORD
          valueFrom:
            secretKeyRef:
              name: db-credentials
              key: password
      volumeMounts:
        - name: secret-volume
          mountPath: /etc/secrets
          readOnly: true
    volumes:
      - name: secret-volume
        secret:
          secretName: db-credentials
```

### **Best Practices:**

- Use ConfigMaps for non-sensitive config
- Enable encryption at rest for secrets
- Use external secret management (HashiCorp Vault)
- Avoid storing secrets in environment variables
- Implement RBAC for secret access

## **SECTION 6: SCALING & RESOURCE MANAGEMENT**

### **Q17: How does Horizontal Pod Autoscaler (HPA) work?**

#### **Answer:**

**HPA** automatically adjusts the number of pod replicas based on metrics.

#### **How HPA Works:**

1. Metrics server collects resource metrics from kubelets
2. HPA controller queries metrics API
3. Compares current metrics with target metrics
4. Calculates desired replica count
5. Updates deployment/statefulset replica count

### **HPA YAML Example:**

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: web-app
  minReplicas: 2
  maxReplicas: 10
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
    - type: Resource
      resource:
        name: memory
        target:
          type: AverageValue
          averageValue: 500Mi
    - type: Pods
      pods:
        metric:
          name: custom_requests_per_second
          target:
            type: AverageValue
            averageValue: "1000"
  behavior:
    scaleDown:
      stabilizationWindowSeconds: 300
  policies:
    - type: Percent
```

```
    value: 50
    periodSeconds: 60
  scaleUp:
    stabilizationWindowSeconds: 0
    policies:
      - type: Percent
        value: 100
        periodSeconds: 30
```

### **Scaling Formula:**

desiredReplicas = ceil(currentMetric / targetMetric) \* currentReplicas

### **Metric Types:**

- **Resource metrics:** CPU, memory
- **Custom metrics:** Application-specific metrics
- **External metrics:** External monitoring systems

### **Best Practices:**

- Set appropriate min/max replicas
- Configure resource requests for accurate scaling
- Use multiple metrics for better decisions
- Monitor scaling events and adjust thresholds

## **Q18: What are Resource Requests and Limits?**

### **Answer:**

Resource requests and limits control CPU and memory allocation.

### **Requests:**

- Minimum guaranteed resources
- Reserved for the pod
- Scheduler uses this for placement
- Pod QoS depends on requests

### **Limits:**

- Maximum resources a container can use
- If exceeded: CPU is throttled, Memory causes OOM kill
- Not guaranteed to be reserved

### **Resource Requests/Limits YAML:**

```
apiVersion: v1
```

```

kind: Pod
metadata:
  name: resource-managed-pod
spec:
  containers:
    - name: app
      image: myapp:1.0
      resources:
        requests:
          cpu: 100m    # 0.1 CPU cores
          memory: 256Mi
        limits:
          cpu: 500m    # 0.5 CPU cores
          memory: 1Gi

```

### **CPU Units:**

- 1 = 1 CPU core = 1000m (millicores)
- 500m = 0.5 CPU cores
- 0.1 = 100m

### **Memory Units:**

- 1Mi = 1 Megabyte
- 1Gi = 1 Gigabyte
- 1Ki = 1 Kilobyte

### **Quality of Service (QoS) Classes:**

#### **1. Guaranteed (Highest Priority)**

- Requests = Limits for both CPU and memory
- Never evicted unless exceeds limits

```

# Guaranteed QoS
resources:
  requests:
    cpu: 100m
    memory: 256Mi
  limits:
    cpu: 100m
    memory: 256Mi

```

#### **2. Burstable (Medium Priority)**

- Requests < Limits
- Evicted if node runs out of resources

```
# Burstable QoS
```

```
resources:  
  requests:  
    cpu: 100m  
    memory: 256Mi  
  limits:  
    cpu: 500m  
    memory: 1Gi
```

### 3. BestEffort (Lowest Priority)

- No requests or limits
- First to be evicted

```
# BestEffort QoS (no resources specified)
```

## SECTION 7: TROUBLESHOOTING & DEBUGGING

### Q19: How do you debug a failing pod?

#### Answer:

Follow systematic debugging steps:

#### Step 1: Check Pod Status

```
kubectl get pods  
kubectl describe pod <pod-name> # Check events and conditions
```

#### Step 2: View Logs

```
kubectl logs <pod-name> # Latest logs  
kubectl logs <pod-name> --previous # Previous crashed container  
kubectl logs <pod-name> -c <container-name> # Specific container  
kubectl logs <pod-name> -f # Follow logs in real-time
```

#### Step 3: Execute Commands Inside Pod

```
kubectl exec -it <pod-name> -- /bin/sh  
kubectl exec -it <pod-name> -- cat /etc/config
```

#### Step 4: Check Resource Availability

```
kubectl describe node <node-name>
kubectl top nodes
kubectl top pods
```

### Step 5: Common Pod Failure States:

State	Cause	Solution
<b>Pending</b>	Insufficient resources, PVC not bound	Add nodes, check PVC/requests
<b>ImagePullBackOff</b>	Image not found, invalid credentials	Check image name, registry credentials
<b>CrashLoopBackOff</b>	Application error, bug	Check logs, fix application
<b>Evicted</b>	Node out of resources	Free resources, adjust limits
<b>NotReady</b>	Readiness probe failing	Check probe configuration, app health

### Debugging Checklist:

1. Pod Events (kubectl describe pod)
2. Container Logs (kubectl logs)
3. Resource Limits & Requests
4. Readiness & Liveness Probes
5. Network Connectivity
6. Volume Mounts & Storage
7. Environment Variables
8. Application Configuration

## Q20: What are Liveness and Readiness Probes?

### Answer:

#### Readiness Probe:

- Checks if pod is ready to receive traffic
- Determines if pod should be part of service endpoints
- Failures mean pod is not ready, traffic not sent

- Used during deployment to avoid sending traffic to initializing pods

### **Liveness Probe:**

- Checks if pod application is still running
- Failures trigger pod restart
- Prevents zombie processes
- Detects deadlocked applications

### **Probe Types:**

1. HTTP
2. TCP Socket
3. Exec Command

### **Probe Configuration YAML:**

```

apiVersion: v1
kind: Pod
metadata:
  name: app-with-probes
spec:
  containers:
  - name: app
    image: myapp:1.0

    # HTTP Readiness Probe
    readinessProbe:
      httpGet:
        path: /health
        port: 8080
        scheme: HTTP
      initialDelaySeconds: 10
      periodSeconds: 5
      timeoutSeconds: 2
      failureThreshold: 3

    # TCP Liveness Probe
    livenessProbe:
      tcpSocket:
        port: 3306
      initialDelaySeconds: 15
      periodSeconds: 10
      timeoutSeconds: 5
      failureThreshold: 3

```

```
# Exec Command Readiness
readinessProbe:
  exec:
    command:
      - /bin/sh
      - -c
      - mysql -h localhost -u root -ppassword -e "SELECT 1"
  initialDelaySeconds: 5
  periodSeconds: 10
```

### **Probe Fields:**

- initialDelaySeconds: Delay before first probe
- periodSeconds: Frequency of probes
- timeoutSeconds: Time to wait for response
- successThreshold: Successes needed to mark as healthy
- failureThreshold: Failures before action taken

### **Best Practices:**

- Implement both readiness and liveness probes
- Use lightweight health check endpoints
- Set appropriate thresholds to avoid false positives
- Don't rely on external service calls in probes

## **SECTION 8: SECURITY & RBAC**

### **Q21: What is Role-Based Access Control (RBAC)?**

#### **Answer:**

**RBAC** controls who can perform what actions on Kubernetes resources.

#### **RBAC Components:**

1. **Role**
  - Namespace-scoped permissions
  - Defines allowed verbs and resources
2. **ClusterRole**
  - Cluster-wide permissions
  - Can be used across namespaces
3. **RoleBinding**
  - Binds Role to user/group/service account (namespace)
4. **ClusterRoleBinding**
  - Binds ClusterRole to user/group/service account (cluster-wide)

## **RBAC YAML Examples:**

```
# ClusterRole for reading pods
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: pod-reader
rules:
- apiGroups: []
  resources: ["pods"]
  verbs: ["get", "list", "watch"]
- apiGroups: []
  resources: ["pods/logs"]
  verbs: ["get"]
```

```
# ClusterRoleBinding
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
  name: read-pods-globally
subjects:
- kind: ServiceAccount
  name: developer
  namespace: default
- kind: User
  name: john@example.com
roleRef:
  kind: ClusterRole
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```

```
# Role for namespace-scoped operations
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: deployment-manager
  namespace: production
rules:
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["create", "get", "update", "patch", "delete"]
```

```

- apiGroups: []
  resources: ["pods", "pods/logs"]
  verbs: ["get", "list", "watch"]

# RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: deploy-manager-binding
  namespace: production
subjects:
- kind: User
  name: alice@example.com
roleRef:
  kind: Role
  name: deployment-manager
  apiGroup: rbac.authorization.k8s.io

```

### **Common Verbs:**

- get, list, watch (Read operations)
- create, update, patch (Write operations)
- delete, deletecollection (Deletion)
- exec, attach (Interactive operations)

### **Service Accounts:**

- Identity for pods
- Automatically mounted as volumes
- Used for pod authentication to API server

```

apiVersion: v1
kind: ServiceAccount
metadata:
  name: app-service-account
  namespace: production

```

## **Q22: What is Pod Security Policy (PSP)?**

### **Answer:**

**Pod Security Policy** defines security standards for pods (being replaced by Pod

Security Admission).

### **PSP Controls:**

- Privileged container execution
- Host network/IPC access
- Host port exposure
- Volume types
- Running as root user
- Capabilities
- AppArmor/SELinux labels

### **PSP Example:**

```
apiVersion: policy/v1beta1
kind: PodSecurityPolicy
metadata:
  name: restricted
spec:
  privileged: false
  allowPrivilegeEscalation: false
  requiredDropCapabilities:
    - ALL
  volumes:
    - 'configMap'
    - 'emptyDir'
    - 'projected'
    - 'secret'
    - 'downwardAPI'
    - 'persistentVolumeClaim'
  hostNetwork: false
  hostIPC: false
  hostPID: false
  runAsUser:
    rule: 'MustRunAsNonRoot'
  seLinux:
    rule: 'MustRunAs'
    seLinuxOptions:
      level: "s0:c123,c456"
  fsGroup:
    rule: 'MustRunAs'
    ranges:
      - min: 1
        max: 65535
```

```
readOnlyRootFilesystem: true
```

### **Pod Security Admission (Kubernetes 1.25+):**

```
apiVersion: v1
kind: Namespace
metadata:
  name: production
  labels:
    pod-security.kubernetes.io/enforce: restricted
    pod-security.kubernetes.io/audit: restricted
    pod-security.kubernetes.io/warn: restricted
```

## **SECTION 9: ADVANCED TOPICS**

### **Q23: What is a Service Mesh and why use it?**

#### **Answer:**

A **Service Mesh** manages service-to-service communication infrastructure layer (Layer 5).

#### **Key Features:**

- 1. Traffic Management**
  - Load balancing algorithms
  - Canary deployments
  - Circuit breaking
  - Retries and timeouts
- 2. Security**
  - Mutual TLS (mTLS) encryption
  - Authentication and authorization
  - Certificate management
- 3. Observability**
  - Distributed tracing
  - Metrics collection
  - Log aggregation
- 4. Resilience**
  - Fault injection testing
  - Traffic mirroring
  - Rate limiting

#### **Popular Service Meshes:**

- **Istio:** Full-featured, complex
- **Linkerd:** Lightweight, high-performance
- **Consul:** Multi-cloud, service networking

#### **Benefits:**

- Transparent traffic control without application changes
- Security policies at infrastructure level
- Better observability of service communication
- Zero trust networking implementation

#### **Challenges:**

- Added complexity and learning curve
- Performance overhead
- Operational burden

### **Q24: What is an Operator?**

#### **Answer:**

A **Kubernetes Operator** extends Kubernetes to manage complex applications.

#### **Components:**

1. **Custom Resource Definition (CRD):** New API resource
2. **Controller:** Business logic that reconciles actual vs desired state
3. **Reconciliation Loop:** Continuous monitoring and correction

#### **Operator Pattern Flow:**

User creates Custom Resource → Operator detects creation → Operator performs setup (create databases, services, configs) → Operator monitors actual state → If mismatch detected, reconcile to desired state → Application managed by operator

#### **Example Operators:**

- **Prometheus Operator:** Manages Prometheus monitoring
- **MySQL Operator:** Manages MySQL database clusters
- **Kafka Operator:** Manages Kafka clusters
- **ArgoCD Operator:** Manages GitOps workflow

#### **Advantages:**

- Automate complex operational tasks
- Encode operational knowledge
- Consistent deployments
- Self-healing applications

### **Q25: What is etcd and its role?**

## **Answer:**

**etcd** is a distributed key-value store that stores Kubernetes cluster state.

## **Role in Kubernetes:**

- Single source of truth for cluster state
- Stores all API objects (pods, services, deployments, etc.)
- Enables cluster consistency
- Provides watch mechanism for state changes

## **Key Characteristics:**

- Strongly consistent
- Highly available (cluster mode)
- Fast (optimized for reads)
- RAFT consensus algorithm
- Transactional support

## **etcd Backup/Restore:**

```
# Backup etcd
ETCDCTL_API=3 etcdctl \
--endpoints=[https://127.0.0.1:2379](https://127.0.0.1:2379) \
--cacert=/etc/kubernetes/pki/etcd/ca.crt \
--cert=/etc/kubernetes/pki/etcd/server.crt \
--key=/etc/kubernetes/pki/etcd/server.key \
snapshot save /backup/etcd-snapshot.db
```

```
# Restore etcd
ETCDCTL_API=3 etcdctl \
--endpoints=[https://127.0.0.1:2379](https://127.0.0.1:2379) \
--cacert=/etc/kubernetes/pki/etcd/ca.crt \
--cert=/etc/kubernetes/pki/etcd/server.crt \
--key=/etc/kubernetes/pki/etcd/server.key \
snapshot restore /backup/etcd-snapshot.db
```

## **Best Practices:**

- Run etcd in cluster mode (3, 5, or 7 nodes)
- Enable encryption at rest
- Implement backup strategy
- Monitor etcd disk I/O and latency
- Use SSD storage for etcd

# **SECTION 10: REAL-WORLD SCENARIOS (MNC Interview Questions)**

## **Q26: How would you migrate a legacy monolithic application to Kubernetes?**

**Answer:**

### **Migration Strategy:**

- **Phase 1: Containerization**
  - Containerize the application (create Dockerfile)
  - Test container locally
  - Push to container registry
  - Ensure no host dependencies
- **Phase 2: Kubernetes Deployment**
  - Create deployment manifest
  - Configure resource requests/limits
  - Set up liveness/readiness probes
  - Configure environment variables
- **Phase 3: Data & State Management**
  - Identify stateful components (databases)
  - Move to StatefulSets if needed
  - Set up persistent volumes
  - Ensure data migration strategy
- **Phase 4: Networking & Service Discovery**
  - Expose via Service (ClusterIP/LoadBalancer)
  - Configure Ingress if needed
  - Update DNS/load balancers
  - Test internal connectivity
- **Phase 5: Monitoring & Logging**
  - Integrate monitoring (Prometheus)
  - Set up centralized logging (ELK, Loki)
  - Create alerts
  - Dashboard creation
- **Phase 6: Zero-Downtime Cutover**
  - Run legacy and Kubernetes versions in parallel
  - Gradual traffic migration
  - Monitor both environments
  - Rollback plan ready
  - Full cutover once stable

### **Challenges & Solutions:**

<b>Challenge</b>	<b>Solution</b>
<b>Database migrations</b>	Use database migration tools, backup strategy
<b>Session management</b>	Implement session store (Redis, database)
<b>Configuration differences</b>	Use ConfigMaps, secrets for environment-specific config
<b>Performance issues</b>	Profile, optimize, use resource requests/limits
<b>Logging complexity</b>	Centralized logging, tracing infrastructure

## **Q27: Design a production-grade Kubernetes cluster for a SaaS platform**

**Answer:**

### **Cluster Architecture:**

Internet → Cloud Load Balancer → Ingress Controller (nginx) → API Services (Multiple Replicas) → Database Services (StatefulSet) / Database (RDS/Managed)

### **Control Plane:**

- Multi-zone deployment (3+ nodes)
- Managed Kubernetes service (EKS, GKE, AKS)
- Auto-backup of etcd
- High availability API server

### **Worker Nodes:**

- Multiple node groups for different workloads
- Auto-scaling enabled
- Regular security patching
- Enough resources for failover

### **Networking:**

- **Ingress for external traffic**
  - Use managed load balancer
  - SSL/TLS termination

- Rate limiting
- DDoS protection
- **Internal networking**
  - Service discovery via CoreDNS
  - Network policies for security

## **Storage:**

- **Application data**
  - Dynamic provisioning via StorageClasses
  - Backup strategy (snapshots)
  - Multi-zone redundancy
- **Database**
  - Managed database service (RDS, CloudSQL)
  - Automated backups
  - Point-in-time recovery
  - Read replicas for scaling

## **Security:**

- **Authentication & Authorization**
  - RBAC for all users
  - Service accounts with minimal permissions
- **Network Security**
  - Network policies to restrict traffic
  - Private cluster (no public IPs for nodes)
  - VPC security groups
- **Data Security**
  - Encryption at rest (etcd, volumes)
  - Encryption in transit (TLS)
  - Secrets management (Vault)
- **Container Security**
  - Image scanning for vulnerabilities
  - Private container registry
  - Pod security policies

## **Monitoring & Logging:**

- **Metrics**
  - Prometheus for collection
  - Grafana for visualization
  - Custom metrics for application
- **Logging**
  - ELK/Loki for centralized logging
  - Fluentd/Filebeat for collection
  - Kibana/Grafana for analysis

- **Tracing**
  - Jaeger for distributed tracing
  - OpenTelemetry SDK in applications
- **Alerting**
  - Alert rules for anomalies
  - Notification channels (PagerDuty, Slack)
  - Runbooks for incidents

### **CI/CD Integration:**

- **Git-based workflow**
  - GitOps for deployments
  - ArgoCD or Flux for reconciliation
  - Automated testing in pipelines
  - Canary/Blue-Green deployments
- **Secrets Management**
  - External secret store (Vault)
  - Automatic rotation
  - RBAC for secret access

### **Disaster Recovery:**

- **Backup Strategy**
  - Regular cluster state backups
  - Database backups (automated)
  - Multi-region setup for failover
  - RTO/RPO defined and tested
- **Scaling**
  - Horizontal Pod Autoscaler
  - Cluster Autoscaler
  - Vertical Pod Autoscaler for fine-tuning

## **Q28: How would you handle scaling issues in a Kubernetes cluster under high traffic?**

### **Answer:**

#### **Identifying Scaling Issues:**

- **Check current resource usage**  
 kubectl top nodes  
 kubectl top pods --all-namespaces
- **Check pod scheduling failures**  
 kubectl get pods --field-selector status.phase=Pending
- **Check HPA status**

```
kubectl get hpa  
kubectl describe hpa <hpa-name>
```

- **Check events for errors**

```
kubectl get events --all-namespaces --sort-by='.lastTimestamp'
```

## Scaling Strategy:

1. **Pod-Level Scaling (HPA)**

```
apiVersion: autoscaling/v2  
kind: HorizontalPodAutoscaler  
metadata:  
  name: api-hpa  
spec:  
  scaleTargetRef:  
    apiVersion: apps/v1  
    kind: Deployment  
    name: api-server  
  minReplicas: 5  
  maxReplicas: 100  
  metrics:  
    - type: Resource  
      resource:  
        name: cpu  
        target:  
          type: Utilization  
          averageUtilization: 70  
    - type: Resource  
      resource:  
        name: memory  
        target:  
          type: Utilization  
          averageUtilization: 80  
  behavior:  
    scaleUp:  
      stabilizationWindowSeconds: 0  
      policies:  
        - type: Percent  
          value: 200  
          periodSeconds: 30  
        selectPolicy: Max  
    scaleDown:  
      stabilizationWindowSeconds: 300
```

```
policies:  
  - type: Percent  
    value: 50  
    periodSeconds: 60
```

## 2. Node-Level Scaling (Cluster Autoscaler)

- Deploy Cluster Autoscaler (Works with AWS EKS, GKE, AKS)

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: cluster-autoscaler  
  namespace: kube-system  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: cluster-autoscaler  
  template:  
    metadata:  
      labels:  
        app: cluster-autoscaler  
    spec:  
      serviceAccountName: cluster-autoscaler  
      containers:  
        - image: k8s.gcr.io/autoscaling/cluster-autoscaler:v1.24.0  
          name: cluster-autoscaler  
          command:  
            - ./cluster-autoscaler  
              - --cloud-provider=aws  
              - --nodes=1:100:asg-name  
              - --skip-nodes-with-local-storage=false
```

## 3. Resource Optimization

- Set appropriate requests and limits

```
resources:  
  requests:  
    cpu: 100m  
    memory: 128Mi  
  limits:  
    cpu: 500m  
    memory: 512Mi
```

- Use resource quotas per namespace

```
apiVersion: v1  
kind: ResourceQuota
```

```
metadata:  
  name: team-quota  
spec:  
  hard:  
    requests.cpu: "100"  
    requests.memory: "200Gi"  
    limits.cpu: "200"  
    limits.memory: "400Gi"
```

#### 4. Connection Pool & Rate Limiting

- o Configure Ingress rate limiting

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
metadata:  
  name: api-ingress  
  annotations:  
    nginx.ingress.kubernetes.io/limit-rps: "1000"  
spec:  
  ingressClassName: nginx  
  rules:  
  - host: api.example.com  
    http:  
      paths:  
      - path: /  
        pathType: Prefix  
      backend:  
        service:  
          name: api-service  
          port:  
            number: 8080
```

#### 5. Cache Layer

- o Deploy Redis cache

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: redis-cache  
spec:  
  replicas: 3  
  selector:  
    matchLabels:  
      app: redis  
  template:  
    metadata:  
      labels:  
        app: redis
```

```
spec:  
  containers:  
    - name: redis  
      image: redis:7.0  
      ports:  
        - containerPort: 6379  
      resources:  
        requests:  
          cpu: 500m  
          memory: 1Gi  
        limits:  
          cpu: 1000m  
          memory: 2Gi
```

### **Monitoring Scaling:**

```
# Watch HPA scaling  
kubectl get hpa -w  
  
# Check scaling events  
kubectl describe hpa api-hpa  
  
# View node autoscaler logs  
kubectl logs -n kube-system -l app=cluster-autoscaler -f  
  
# Metrics monitoring  
kubectl get --raw /apis/metrics.k8s.io/v1beta1/nodes
```

## **Q29: How would you handle a Pod that crashes repeatedly?**

### **Answer:**

#### **Step 1: Identify the Issue**

```
kubectl describe pod <pod-name> # Check events and conditions  
kubectl logs <pod-name> --previous # View logs from previous crash  
kubectl logs <pod-name> -f # Follow current logs
```

### **Common Causes & Solutions:**

Issue	Symptoms	Solution

<b>Application Bug</b>	Container exits immediately	Review logs, fix code, redeploy
<b>Memory Leak</b>	OOMKilled	Increase memory limit, profile app
<b>Infinite Loop</b>	CPU spike then restart	Debug application code
<b>External Dependency</b>	Cannot connect to database	Check connectivity, ensure dependency is running
<b>Configuration Error</b>	Application fails to start	Validate ConfigMaps, Secrets, environment variables
<b>Image Issue</b>	ImagePullBackOff	Check image name, registry credentials
<b>Resource Limits</b>	Insufficient CPU/memory	Increase limits, optimize application

### **CrashLoopBackOff:**

- Kubernetes auto-restart with exponential backoff
- Delays: 10s, 20s, 40s, 80s, 160s, 300s (max)
- To prevent restart:  
spec:  
  restartPolicy: Never # Or OnFailure
- To debug without restart:  
  kubectl run debug-pod --image alpine --restart=Never -- sleep 3600  
  kubectl exec -it debug-pod -- sh

### **Debugging Steps:**

1. **Check events**  
  kubectl describe pod <pod-name>
2. **View application logs**  
  kubectl logs <pod-name>  
  kubectl logs <pod-name> --previous

### 3. **Check resource usage**

```
kubectl describe pod <pod-name> | grep -A 5 "Limits\|Requests"
```

### 4. **Check environment and configs**

```
kubectl exec <pod-name> -- env
```

```
kubectl exec <pod-name> -- cat /etc/config/app.conf
```

### 5. **Test connectivity**

```
kubectl exec <pod-name> -- ping database-service
```

```
kubectl exec <pod-name> -- curl http://api-service:8080/health
```

### 6. **Check readiness probe**

```
kubectl get pod <pod-name> -o yaml | grep -A 10 "readinessProbe"
```

### 7. **Test locally**

```
docker run --rm myapp:latest
```

## **Prevention Strategies:**

- **Implement readiness probes**

```
readinessProbe:
```

```
  httpGet:
```

```
    path: /health
```

```
    port: 8080
```

```
  initialDelaySeconds: 10
```

```
  periodSeconds: 5
```

```
  failureThreshold: 3
```

- **Set resource limits**

```
resources:
```

```
  requests:
```

```
    cpu: 100m
```

```
    memory: 256Mi
```

```
  limits:
```

```
    cpu: 500m
```

```
    memory: 1Gi
```

- **Use init containers for setup**

```
initContainers:
```

```
- name: setup
```

```
  image: busybox
```

```
  command: ['sh', '-c', 'until nslookup database; do echo waiting; sleep 2; done']
```

```
done']
```

- **Add startup probe (K8s 1.18+)**

```
startupProbe:
  httpGet:
    path: /startup
    port: 8080
  initialDelaySeconds: 0
  periodSeconds: 10
  failureThreshold: 30
```

## **Q30: Design a disaster recovery strategy for a Kubernetes cluster**

**Answer:**

**RTO/RPO Definition:**

- **RTO (Recovery Time Objective):** Maximum acceptable downtime (target: 15-30 minutes)
- **RPO (Recovery Point Objective):** Maximum data loss acceptable (target: 5-15 minutes)

**Multi-Layer DR Strategy:**

1. **Cluster-Level Backup:**

- **Backup etcd (cluster state)**

```
ETCDCTL_API=3 etcdctl snapshot save etcd-backup-$(date +%s).db
```

- **Script for automated backups**

```
#!/bin/bash
BACKUP_DIR="/backups/etcd"
mkdir -p $BACKUP_DIR
```

```
ETCDCTL_API=3 etcdctl \
--endpoints=[https://127.0.0.1:2379](https://127.0.0.1:2379) \
--cacert=/etc/kubernetes/pki/etcd/ca.crt \
--cert=/etc/kubernetes/pki/etcd/server.crt \
--key=/etc/kubernetes/pki/etcd/server.key \
snapshot save $BACKUP_DIR/etcd-$(date +%Y%m%d-%H%M%S).db
```

```
# Cleanup old backups (keep last 7 days)
find $BACKUP_DIR -type f -mtime +7 -delete
```

2. **Application-Level Backup (Velero):**

- **Backup volumes**

```

apiVersion: v1
kind: CronJob
metadata:
  name: volume-backup
spec:
  schedule: "*/5 * * * *" # Every 5 minutes
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: backup
              image: velero:latest
              command:
                - velero
                - backup
                - --create
                - --include-namespaces=*
        restartPolicy: OnFailure

```

### 3. Multi-Region Setup:

- **Region 1 (Primary)**
  - Cluster A (Active)
  - Database (Primary)
  - Load Balancer
- **Region 2 (Standby)**
  - Cluster B (Standby)
  - Database (Replica)
  - Load Balancer
- **Global DNS (Route53, CloudDNS)**
  - Route 90% to Region 1
  - Route 10% to Region 2 (canary)

### 4. Database Replication:

- **Primary-Standby replication**
  - Primary DB: Write operations, Replication to Standby
  - Standby DB: Read-only replica, Can be promoted to primary

### 5. GitOps for Infrastructure Recovery:

- Store all configurations in Git
- Use ArgoCD to reconcile cluster state

```

apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:

```

```

name: app-deployment
spec:
  project: default
  source:
    repoURL:
      [https://github.com/company/k8s-configs](https://github.com/company/k8s-configs)
    targetRevision: main
    path: production/
  destination:
    server: [https://kubernetes.default.svc](https://kubernetes.default.svc)
    namespace: production
  syncPolicy:
    automated:
      prune: true
      selfHeal: true

```

## 6. Testing DR:

- **Regular DR drills**
  1. Backup cluster to secondary
  2. Restore from backup to test cluster
  3. Verify data consistency
  4. Validate application functionality
  5. Measure RTO
  6. Document findings and improvements
- **Automated DR test script**

```

#!/bin/bash
# Create test cluster from backup
# Run smoke tests
# Measure recovery time
# Generate report

```

## 7. Monitoring & Alerting:

- **Monitor backup health**

```

alert: BackupFailed
expr: backup_success{namespace="velero"} == 0
for: 1h
annotations:
  summary: "Velero backup failed"

```
- **Monitor replication lag**
  - alert: DatabaseReplicationLag
 

```

expr: replication_lag_seconds > 60
annotations:

```

summary: "Database replication lagging"

- **Monitor cluster health**  
- alert: KubernetesNodeDown  
  expr: up{job="kubernetes-nodes"} == 0  
  for: 5m

## KEY TAKEAWAYS

### For MNC Interviews:

1. Master core concepts (Pods, Deployments, Services, Storage)
2. Understand troubleshooting and debugging methodologies
3. Know security best practices (RBAC, Network Policies, Secrets Management)
4. Be prepared for real-world scenarios (migrations, scaling, DR)
5. Demonstrate hands-on experience with kubectl and manifests
6. Understand production considerations (monitoring, logging, HA)
7. Know different deployment strategies (Rolling, Blue-Green, Canary)
8. Understand networking and service mesh concepts
9. Be familiar with ecosystem tools (Helm, Prometheus, Loki, Istio)
10. Practice troubleshooting by actually deploying and breaking things

### Preparation Tips:

- Set up a local Kubernetes cluster (Minikube, Docker Desktop)
- Deploy real applications and troubleshoot issues

**Thank you. Lets Connect**

[www.linkedin.com/in/bhooshan-pattanashetti](https://www.linkedin.com/in/bhooshan-pattanashetti)